

▼ Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
from google.colab import drive
drive.flush_and_unmount()
drive.mount("/content/gdrive", force_remount=True)

Drive not mounted, so nothing to flush and unmount.
Mounted at /content/gdrive

import os
os.chdir("/content/gdrive/MyDrive/CS444/assignment4_materials/")

import zipfile
train_data_zip = zipfile.ZipFile("/content/gdrive/MyDrive/CS444/assignment4_materials/cats.zip", "r") #Opens the tar file in read mode
train_data_zip.extractall("/tmp") #Extracts the files into the /tmp folder
train_data_zip.close()

import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

from gan.train import train

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

▼ GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

▼ GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log(1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
from gan.losses import discriminator_loss, generator_loss
```

▼ Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
from gan.losses import ls_discriminator_loss, ls_generator_loss
```

▼ GAN model architecture

TODO: Implement the `Discriminator` and `Generator` networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution).

This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with `in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1`
- batch norm
- transpose convolution with `in_channels=1024, out_channels=512, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=512, out_channels=256, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=256, out_channels=128, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=128, out_channels=3, kernel=4, stride=2`

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a `3x64x64` tensor for each sample (equal dimensions to the images from the dataset).

```
from gan.models import Discriminator, Generator
```

▼ Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

TODO: Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
batch_size = 32
imsize = 64
cat_root = '/tmp/cats'

cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imszie)),
    transforms.RandomCrop(imszie),
]))

cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)
```

▼ Visualize dataset

```
from gan.utils import show_images

imgs = next(cat_loader_train.__iter__())[0].numpy().squeeze() #cat_loader_train.__iter__().next()[0].numpy().squeeze()
show_images(imgs, color=True)

/usr/local/lib/python3.9/dist-packages/torchvision/transforms/functional.py:1603: UserWarning: The default value of the anti-aliasing argument has changed, and now always applies a Gaussian kernel. You can get the previous behavior by setting anti_aliasing=False.
  warnings.warn("The default value of the anti-
```



▼ Training

TODO: Fill in the training loop in `gan/train.py`.

```
NOISE_DIM = 100
NUM_EPOCHS = 50
learning_rate = 0.001
```

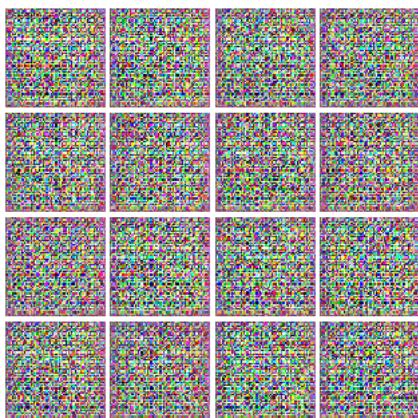
▼ Train GAN

```
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

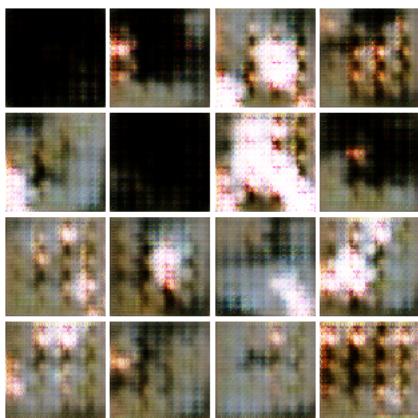
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

# original gan
train(D, G, D_optimizer, G_optimizer, discriminator_loss,
      generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 1.404, G:2.805



Iter: 250, D: 0.9422, G:1.7

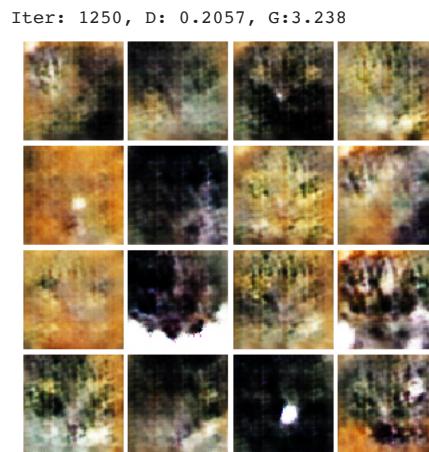
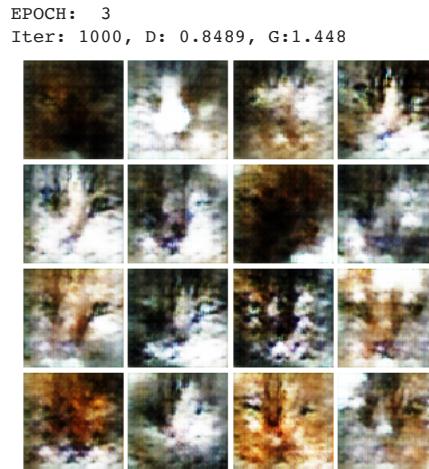


EPOCH: 2
Iter: 500, D: 1.711, G:1.24



Iter: 750, D: 0.9157, G:1.882





EPOCH: 5

Iter: 2000, D: 1.098, G:5.066



Iter: 2250, D: 0.4236, G:2.732



EPOCH: 6

Iter: 2500, D: 0.733, G:3.896



Iter: 2750, D: 1.07, G:1.73





EPOCH: 7
Iter: 3000, D: 0.8408, G:4.21



Iter: 3250, D: 0.2734, G:4.345



EPOCH: 8
Iter: 3500, D: 0.86, G:2.214



Iter: 3750, D: 0.3455, G:3.434



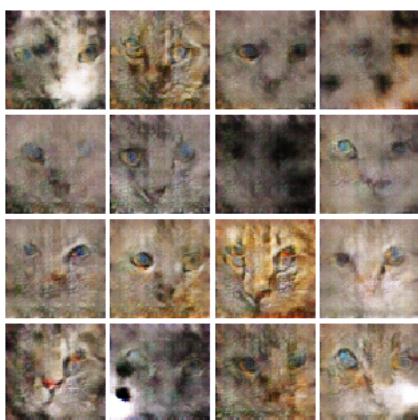




EPOCH: 11
Iter: 5000, D: 0.03565, G:5.882



Iter: 5250, D: 0.4591, G:20.85



EPOCH: 12
Iter: 5500, D: 0.2402, G:4.629



Iter: 5750, D: 0.1976, G:4.249





EPOCH: 13
Iter: 6000, D: 0.3076, G:5.021



Iter: 6250, D: 1.181, G:11.03



EPOCH: 14
Iter: 6500, D: 0.07048, G:4.843



Iter: 6750, D: 0.2046, G:3.685





EPOCH: 15

Iter: 7000, D: 0.1427, G:5.238



Iter: 7250, D: 0.6598, G:6.221



EPOCH: 16

Iter: 7500, D: 0.2301, G:2.761



Iter: 7750, D: 0.08153, G:6.878





EPOCH: 17
Iter: 8000, D: 0.2159, G:5.846



Iter: 8250, D: 0.1186, G:5.993



EPOCH: 18
Iter: 8500, D: 0.03187, G:6.758



Iter: 8750, D: 0.05047, G:6.51





EPOCH: 19

Iter: 9000, D: 0.2792, G:6.478



Iter: 9250, D: 0.06373, G:7.873

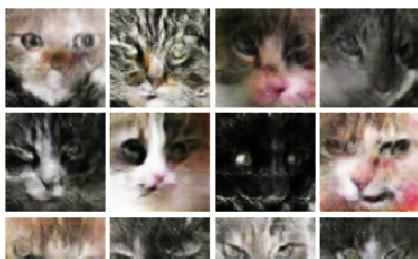


EPOCH: 20

Iter: 9500, D: 0.0853, G:3.703



Iter: 9750, D: 0.2928, G:9.973





EPOCH: 21
Iter: 10000, D: 0.02043, G:8.167



Iter: 10250, D: 0.07849, G:9.583



EPOCH: 22
Iter: 10500, D: 0.07928, G:5.885



Iter: 10750, D: 0.02407, G:6.817





EPOCH: 23
Iter: 11000, D: 0.04304, G:5.598



Iter: 11250, D: 0.006365, G:7.721



EPOCH: 24
Iter: 11500, D: 0.0651, G:7.385



Iter: 11750, D: 0.5371, G:5.951





EPOCH: 25

Iter: 12000, D: 0.05112, G:6.963



Iter: 12250, D: 0.04958, G:7.433



EPOCH: 26

Iter: 12500, D: 0.07588, G:6.511



Iter: 12750, D: 4.791, G:5.498





EPOCH: 27
Iter: 13000, D: 0.01247, G:8.056



Iter: 13250, D: 0.04602, G:6.558



EPOCH: 28
Iter: 13500, D: 0.03023, G:8.067



Iter: 13750, D: 0.07723, G:5.593





EPOCH: 29

Iter: 14000, D: 0.2474, G:5.548



Iter: 14250, D: 0.1227, G:5.541



EPOCH: 30

Iter: 14500, D: 0.00997, G:7.735



Iter: 14750, D: 0.009817, G:7.568





EPOCH: 31
Iter: 15000, D: 0.08608, G:7.984



Iter: 15250, D: 0.01281, G:7.775



EPOCH: 32
Iter: 15500, D: 0.1205, G:5.049



EPOCH: 33
Iter: 15750, D: 0.02738, G:8.327





Iter: 16000, D: 0.1542, G:8.123



EPOCH: 34

Iter: 16250, D: 0.01862, G:7.812



Iter: 16500, D: 0.0076, G:8.734



EPOCH: 35

Iter: 16750, D: 0.02944, G:5.054





Iter: 17000, D: 0.008898, G:12.08

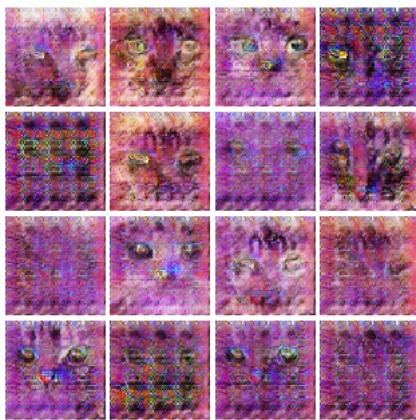


EPOCH: 36

Iter: 17250, D: 0.01438, G:7.966



Iter: 17500, D: 0.0002601, G:12.87



EPOCH: 37

Iter: 17750, D: 0.09022, G:7.71





Iter: 18000, D: 0.05768, G:5.673



EPOCH: 38

Iter: 18250, D: 0.1524, G:6.66



Iter: 18500, D: 0.7236, G:12.54



EPOCH: 39

Iter: 18750, D: 0.01523, G:6.444





Iter: 19000, D: 0.008833, G:7.177



EPOCH: 40

Iter: 19250, D: 0.01903, G:4.792



Iter: 19500, D: 0.03094, G:6.655



EPOCH: 41

Iter: 19750, D: 0.1572, G:6.271





Iter: 20000, D: 0.02999, G:7.638



EPOCH: 42

Iter: 20250, D: 0.02467, G:7.308



Iter: 20500, D: 0.01974, G:8.172



EPOCH: 43

Iter: 20750, D: 0.02644, G:6.272





Iter: 21000, D: 0.03563, G:2.99



EPOCH: 44

Iter: 21250, D: 0.3641, G:4.236



Iter: 21500, D: 0.01043, G:8.536



EPOCH: 45

Iter: 21750, D: 0.1796, G:8.365





Iter: 22000, D: 0.01198, G:10.12



EPOCH: 46

Iter: 22250, D: 0.1503, G:5.476



Iter: 22500, D: 0.009958, G:5.899



EPOCH: 47

Iter: 22750, D: 0.009393, G:6.668



Iter: 23000, D: 0.07664, G:7.575



EPOCH: 48

Iter: 23250, D: 0.0244, G:7.314



Iter: 23500, D: 4.589, G:16.19



EPOCH: 49

Iter: 23750, D: 0.0257, G:8.201



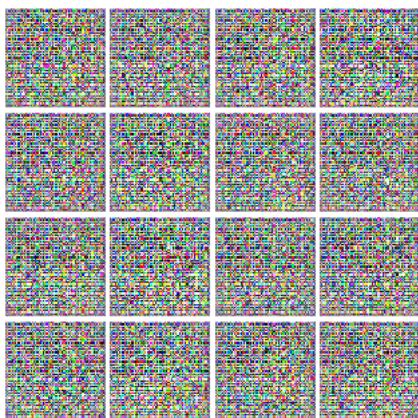
▼ Train LS-GAN

```
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

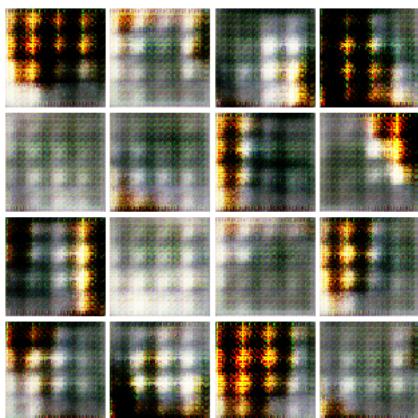
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

# ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

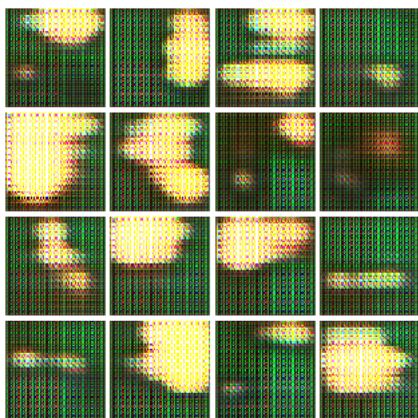
EPOCH: 1
Iter: 0, D: 0.5973, G:35.1



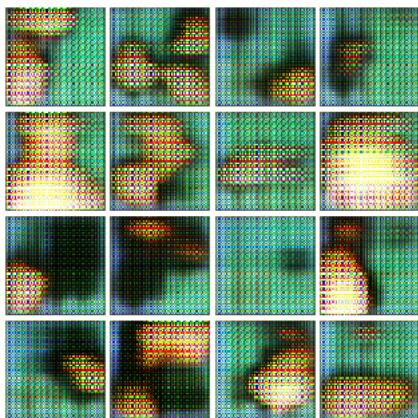
Iter: 250, D: 0.2108, G:0.2752



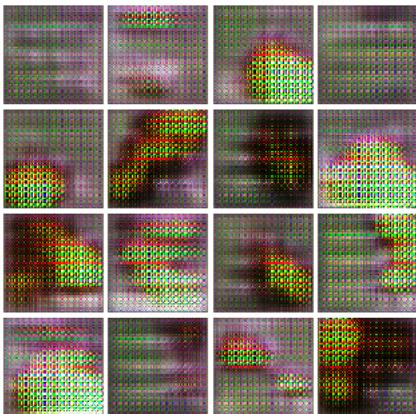
EPOCH: 2
Iter: 500, D: 0.1851, G:0.2248



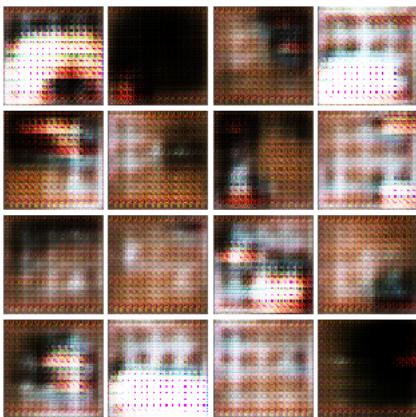
Iter: 750, D: 0.265, G:0.1572



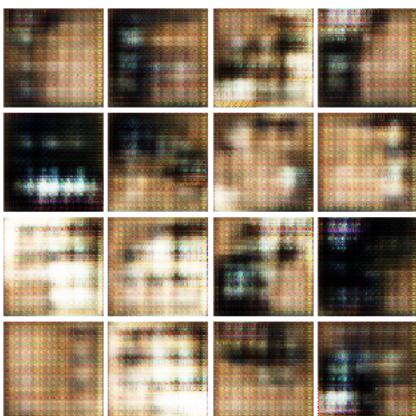
EPOCH: 3
Iter: 1000, D: 0.1566, G:0.1761



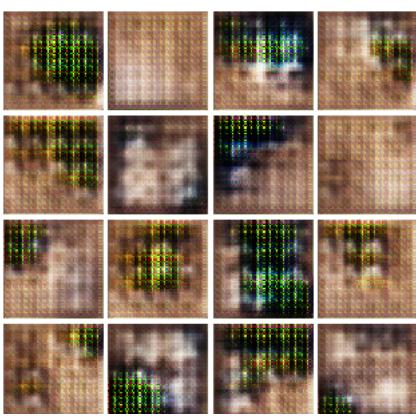
Iter: 1250, D: 0.2173, G:0.1472



EPOCH: 4
Iter: 1500, D: 0.2639, G:0.1442

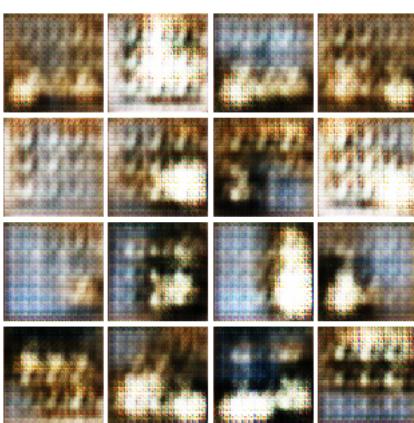


Iter: 1750, D: 0.7418, G:0.2387

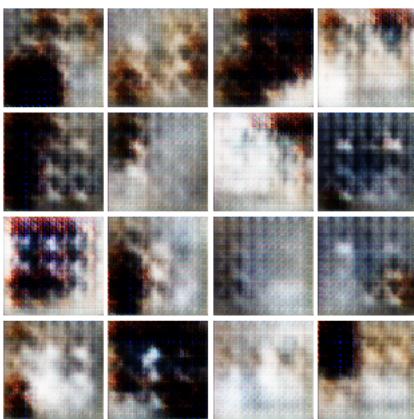


EPOCH: 5

Iter: 2000, D: 0.1871, G:0.3416



Iter: 2250, D: 0.2343, G:0.3687

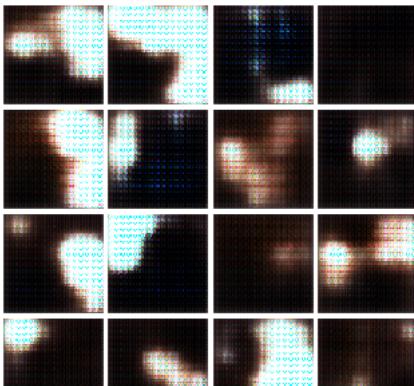


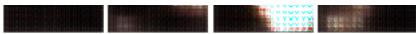
EPOCH: 6

Iter: 2500, D: 0.326, G:0.3762

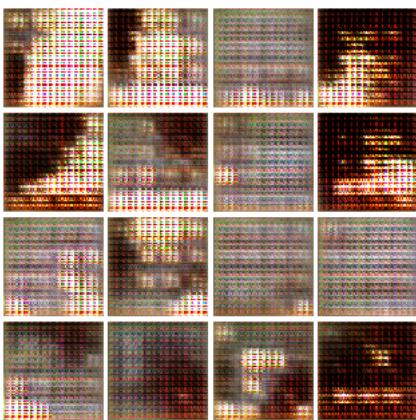


Iter: 2750, D: 0.1983, G:0.3028

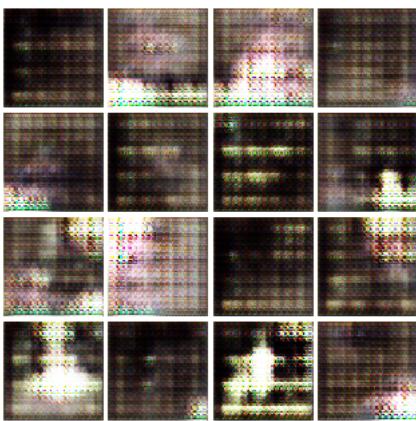




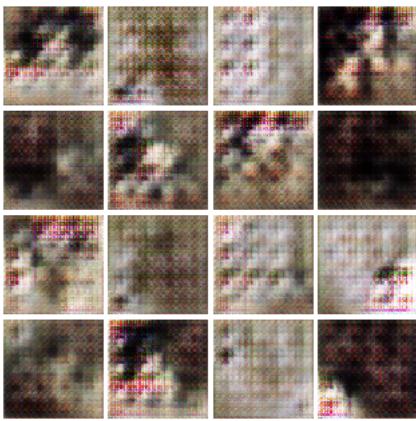
EPOCH: 7
Iter: 3000, D: 0.1883, G: 0.2189



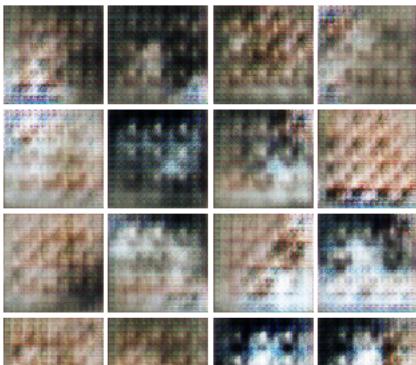
Iter: 3250, D: 0.2432, G: 0.1588



EPOCH: 8
Iter: 3500, D: 0.2706, G: 0.1392



Iter: 3750, D: 0.2559, G: 0.1653





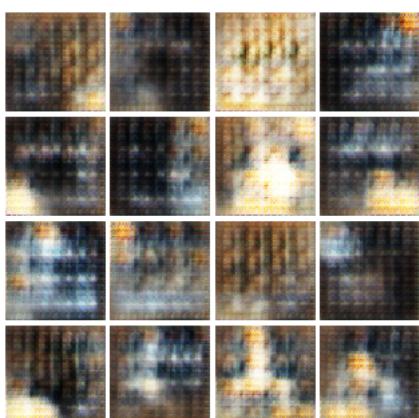
EPOCH: 9
Iter: 4000, D: 0.2316, G:0.1162



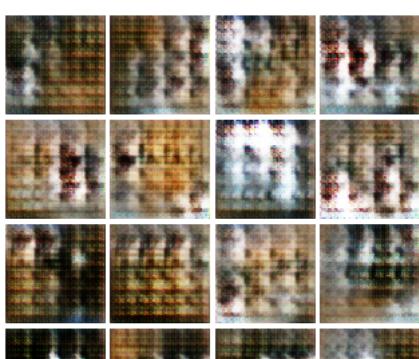
Iter: 4250, D: 0.5388, G:0.1739



EPOCH: 10
Iter: 4500, D: 0.2408, G:0.1288

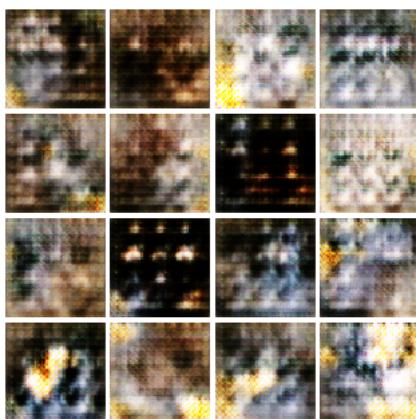


Iter: 4750, D: 0.2778, G:0.2098

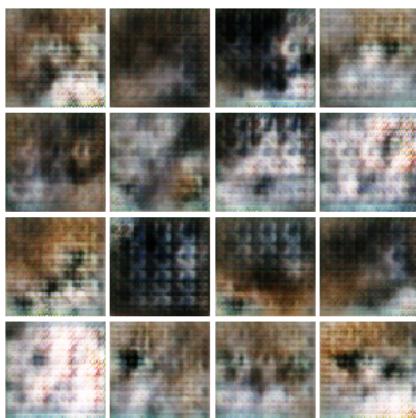




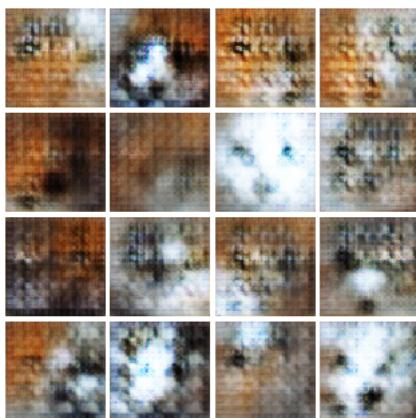
EPOCH: 11
Iter: 5000, D: 0.2643, G: 0.1861



Iter: 5250, D: 0.2309, G: 0.1592



EPOCH: 12
Iter: 5500, D: 0.2631, G: 0.2627

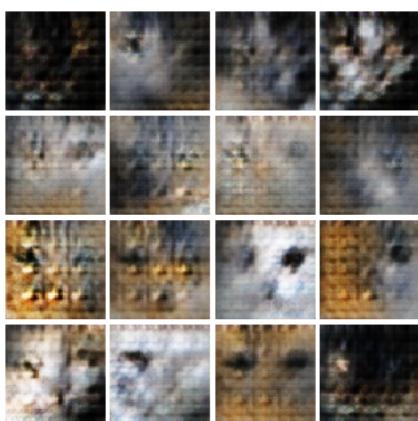


Iter: 5750, D: 0.2272, G: 0.18

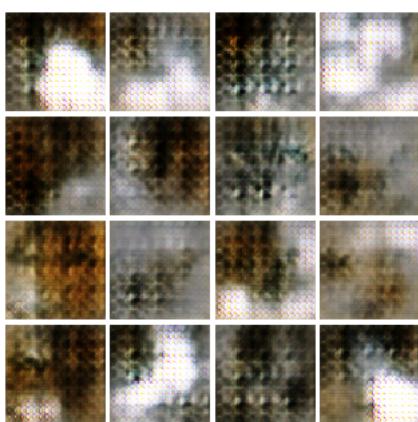




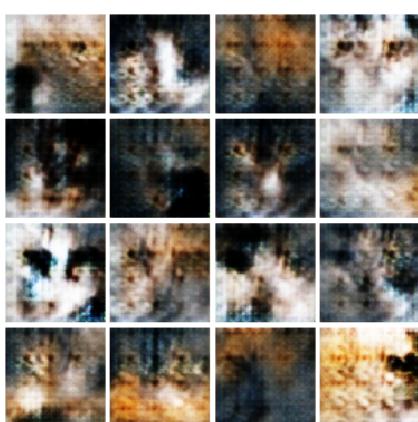
EPOCH: 13
Iter: 6000, D: 0.2555, G: 0.1714



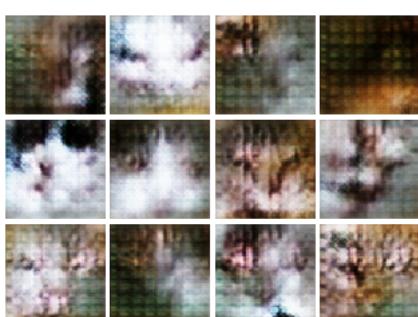
Iter: 6250, D: 0.2312, G: 0.205



EPOCH: 14
Iter: 6500, D: 0.2717, G: 0.1736



Iter: 6750, D: 0.2445, G: 0.2322





EPOCH: 15
Iter: 7000, D: 0.2032, G:0.2348



Iter: 7250, D: 0.2839, G:0.1844



EPOCH: 16
Iter: 7500, D: 0.2822, G:0.2



Iter: 7750, D: 0.241, G:0.1708





EPOCH: 17

Iter: 8000, D: 0.1979, G: 0.1908

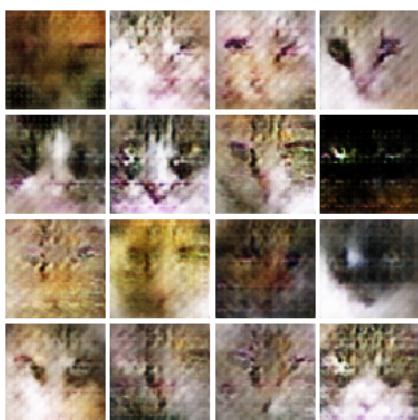


Iter: 8250, D: 0.2508, G: 0.244



EPOCH: 18

Iter: 8500, D: 0.206, G: 0.3474



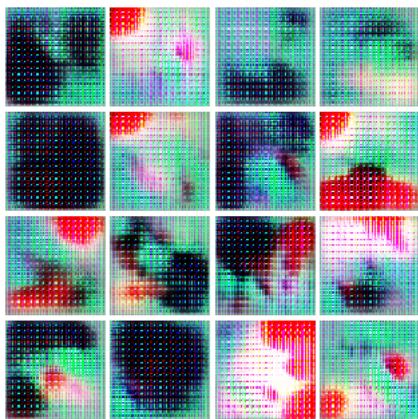
Iter: 8750, D: 0.2228, G: 0.3239



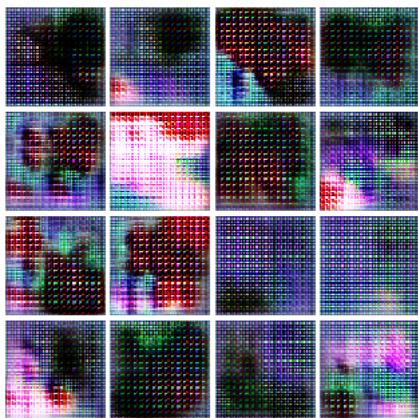


EPOCH: 19

Iter: 9000, D: 0.0934, G: 0.285

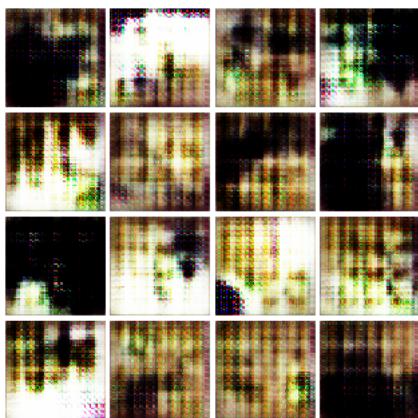


Iter: 9250, D: 0.06444, G: 0.3113



EPOCH: 20

Iter: 9500, D: 0.05643, G: 0.4045

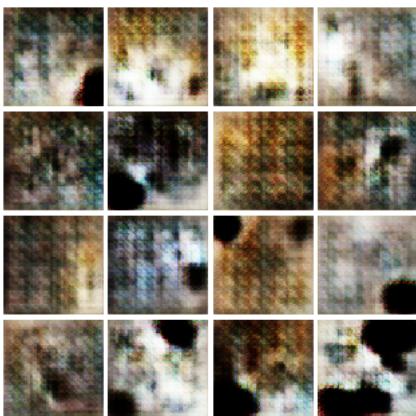


Iter: 9750, D: 0.09508, G: 0.3949

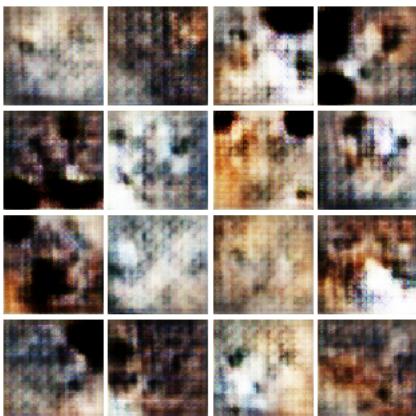




EPOCH: 21
Iter: 10000, D: 0.1827, G:0.176



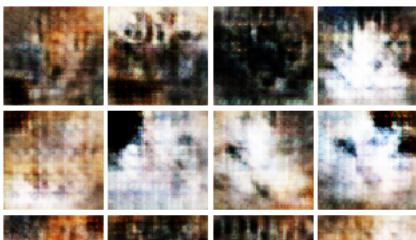
Iter: 10250, D: 0.2029, G:0.1859



EPOCH: 22
Iter: 10500, D: 0.198, G:0.2245



Iter: 10750, D: 0.1406, G:0.3166





EPOCH: 23
Iter: 11000, D: 0.2578, G:0.21



Iter: 11250, D: 0.2517, G:0.2031



EPOCH: 24
Iter: 11500, D: 0.2426, G:0.179



Iter: 11750, D: 0.2328, G:0.1269





EPOCH: 25

Iter: 12000, D: 0.2919, G:0.3037



Iter: 12250, D: 0.1978, G:0.2729

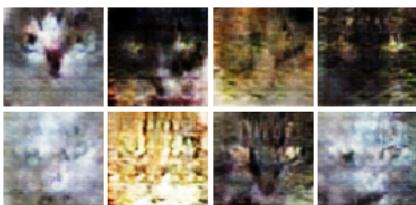


EPOCH: 26

Iter: 12500, D: 0.2646, G:0.1056



Iter: 12750, D: 0.2752, G:0.3823





EPOCH: 27
Iter: 13000, D: 1.101, G:0.3729



Iter: 13250, D: 0.1774, G:0.177



EPOCH: 28
Iter: 13500, D: 0.2413, G:0.1693



Iter: 13750, D: 0.1414, G:0.4044





EPOCH: 29

Iter: 14000, D: 0.2507, G:0.4903



Iter: 14250, D: 0.1265, G:0.2516

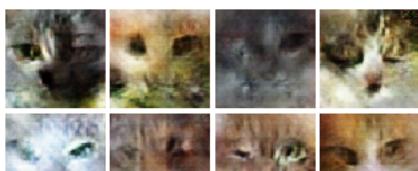


EPOCH: 30

Iter: 14500, D: 0.2378, G:0.1357



Iter: 14750, D: 0.2389, G:0.159





EPOCH: 31
Iter: 15000, D: 0.2292, G:0.1883



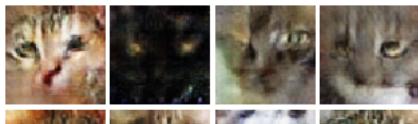
Iter: 15250, D: 0.2006, G:0.352



EPOCH: 32
Iter: 15500, D: 0.1559, G:0.6651



EPOCH: 33
Iter: 15750, D: 0.2461, G:0.3297





Iter: 16000, D: 0.3649, G:0.2076



EPOCH: 34

Iter: 16250, D: 0.1084, G:0.3659



Iter: 16500, D: 0.2306, G:0.1273



EPOCH: 35

Iter: 16750, D: 0.3029, G:0.1293





Iter: 17000, D: 0.2617, G:0.2347



EPOCH: 36

Iter: 17250, D: 0.2708, G:0.1142



Iter: 17500, D: 0.4334, G:0.1332



EPOCH: 37

Iter: 17750, D: 0.2049, G:0.1132





Iter: 18000, D: 0.2841, G: 0.1399



EPOCH: 38

Iter: 18250, D: 0.1827, G: 0.225



Iter: 18500, D: 0.2338, G: 0.1479



EPOCH: 39

Iter: 18750, D: 0.2027, G: 0.2898





Iter: 19000, D: 0.2607, G:0.1212



EPOCH: 40

Iter: 19250, D: 0.1607, G:0.2037



Iter: 19500, D: 0.2503, G:0.2848



EPOCH: 41

Iter: 19750, D: 0.1847, G:0.2524





Iter: 20000, D: 0.2066, G:0.1902



EPOCH: 42

Iter: 20250, D: 0.2422, G:0.1601



Iter: 20500, D: 0.2393, G:0.1518



EPOCH: 43

Iter: 20750, D: 0.2737, G:0.4021





Iter: 21000, D: 0.1811, G:0.3463



EPOCH: 44

Iter: 21250, D: 0.1577, G:0.3741



Iter: 21500, D: 0.21, G:0.225



EPOCH: 45

Iter: 21750, D: 0.1994, G:0.9821





Iter: 22000, D: 0.1457, G:0.2982



EPOCH: 46

Iter: 22250, D: 0.1709, G:0.2385



Iter: 22500, D: 0.1451, G:0.2722



EPOCH: 47

Iter: 22750, D: 0.121, G:0.4987



Iter: 23000, D: 0.1642, G:0.208



EPOCH: 48

Iter: 23250, D: 0.141, G:0.2614



Iter: 23500, D: 0.1876, G:0.266



EPOCH: 49

Iter: 23750, D: 0.1435, G:0.3286