

▼ GAN Debugging Notebook

This notebook is provided to help you debug your code. We provide you with small discriminator and generator networks that you can train on the MNIST dataset. This small GAN can be trained quickly on MNIST and will help you verify that your loss functions and training code is correct.

```
from google.colab import drive
drive.flush_and_unmount()
drive.mount("/content/gdrive",force_remount=True)

Mounted at /content/gdrive

import os
os.chdir("/content/gdrive/MyDrive/assignment4_materials/")

import torch
import torch.nn as nn
from torchvision import datasets
from torchvision import transforms
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder

import numpy as np

import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

Show hidden output

from gan.train import train
from gan.utils import sample_noise, show_images, deprocess_img, preprocess_img
from gan.losses import discriminator_loss, generator_loss, ls_discriminator_loss, ls_generator_loss

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

▼ MNIST Dataset

```
NOISE_DIM = 100
batch_size = 128

mnist = datasets.MNIST('./MNIST_data', train=True, download=True,
                      transform=transforms.ToTensor())
loader_train = DataLoader(mnist, batch_size=batch_size, drop_last=True)

imgs = next(loader_train.__iter__())[0].view(batch_size, 784).numpy().squeeze() #loader_train.__iter__().next()[0].view(batch_size, 784).numpy().squeeze()
show_images(imgs)

Show hidden output
```

▼ Discriminator and Generator

```
class Flatten(nn.Module):
    def forward(self, x):
        #print(x.shape)
        #N, C, H, W = x.size() # read in N, C, H, W
```

```
return x.view(x.shape[0], -1) # "flatten" the C * H * W values into a single vector per image
```

```
def discriminator():
    """
    Initialize and return a simple discriminator model.
    """
    model = torch.nn.Sequential( Flatten(),
                                torch.nn.Linear(784, 256),
                                torch.nn.LeakyReLU(),
                                torch.nn.Linear(256, 256),
                                torch.nn.LeakyReLU(),
                                torch.nn.Linear(256, 1)
                                )
    return model

def generator(noise_dim=NOISE_DIM):
    """
    Initialize and return a simple generator model.
    """

    model = nn.Sequential(
        Flatten(),
        torch.nn.Linear(noise_dim, 1024),
        torch.nn.ReLU(),
        torch.nn.Linear(1024, 1024),
        torch.nn.ReLU(),
        torch.nn.Linear(1024, 784),
        torch.nn.Tanh()
    )

    return model
```

Test to make sure the number of parameters in the generator is correct:

▼ Train

The simple model provided will train on MNIST in only a few minutes. You should expect results that resemble the following if your loss function and training loop implementations are correct:




```
# original GAN
D = discriminator().to(device)
G = generator().to(device)


D_optimizer = torch.optim.Adam(D.parameters(), lr=1e-3, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=1e-3, betas = (0.5, 0.999))

train(D, G, D_optimizer, G_optimizer, discriminator_loss, generator_loss, train_loader=loader_train, num_epochs=10, device=device)
```


EPOCH: 1
Iter: 0, D: 0.6632, G:0.7255




Iter: 250, D: 0.6523, G:1.218



EPOCH: 2
Iter: 500, D: 0.3751, G:1.334




Iter: 750, D: 0.5195, G:0.8959




EPOCH: 3
Iter: 1000, D: 0.467, G:1.722




Iter: 1250, D: 0.5017, G:1.378



EPOCH: 4
Iter: 1500, D: 0.5831, G:1.069



Iter: 1750, D: 0.5811, G:1.017



EPOCH: 5

Iter: 2000, D: 0.6385, G:0.7514



Iter: 2250, D: 0.6345, G:0.9456




EPOCH: 6

Iter: 2500, D: 0.6589, G:0.9116




Iter: 2750, D: 0.6971, G:0.8735






EPOCH: 7

Iter: 3000, D: 0.6435, G:0.8274



Iter: 3250, D: 0.6186, G:0.8112




EPOCH: 8

Iter: 3500, D: 0.6491, G:0.8476



EPOCH: 9

Iter: 3750, D: 0.6364, G:0.7904





Iter: 4000, D: 0.7122, G:0.7916



EPOCH: 10

Iter: 4250, D: 0.6272, G:0.8375



```
# LSGAN
```

```
D_LS = discriminator().to(device)
```


```
G_LS = generator().to(device)
```

```
D_LS_optimizer = torch.optim.Adam(D_LS.parameters(), lr=1e-3, betas = (0.5, 0.999))
```


```
G_LS_optimizer = torch.optim.Adam(G_LS.parameters(), lr=1e-3, betas = (0.5, 0.999))
```

```
train(D_LS, G_LS, D_LS_optimizer, G_LS_optimizer, ls_discriminator_loss, ls_generator_loss, train_loader=loader_train, num_epochs=
```


EPOCH: 1
Iter: 0, D: 0.3805, G:0.4103




Iter: 250, D: 0.2892, G:0.2648




EPOCH: 2
Iter: 500, D: 0.2336, G:0.02312




Iter: 750, D: 0.1238, G:0.4017



EPOCH: 3
Iter: 1000, D: 0.1388, G:0.2856



Iter: 1250, D: 0.1644, G:0.2095



EPOCH: 4
Iter: 1500, D: 0.15, G:0.2473



Iter: 1750, D: 0.1815, G:0.2463



EPOCH: 5

Iter: 2000, D: 0.1951, G:0.1983



Iter: 2250, D: 0.2001, G:0.1794



EPOCH: 6

Iter: 2500, D: 0.206, G:0.2107




Iter: 2750, D: 0.2285, G:0.1876





EPOCH: 7

Iter: 3000, D: 0.2037, G:0.1946




Iter: 3250, D: 0.2037, G:0.2381




EPOCH: 8

Iter: 3500, D: 0.2325, G:0.1569



EPOCH: 9

Iter: 3750, D: 0.2084, G:0.1882






Iter: 4000, D: 0.2241, G:0.1662



EPOCH: 10

Iter: 4250, D: 0.2047, G:0.2019



Iter: 4500, D: 0.2305, G:0.1599

